



All 7 algorithms · 10 new seeds · entropy $\beta = 0.1$

Original validation reward



Logarithmic distance to the upper bound



— Online QCV — RLOO — REINFORCE — GRPO
— Online Q + row delta — RELAX — OTB

Per-epoch validation mean, not best-so-far. Shading: pointwise 95% t intervals across 10 seeds, transformed identically. Log transform is applied after averaging rewards; upper bound 1 is not assumed attainable. No smoothing or fitted scale.